



Apache Spark on Amazon EMR

Pooja Chikkala, Solutions Architect

04/20/2020

Agenda


Introduction to Apache Spark
Spark on Amazon EMR
Tuning best practices

FAST




RICH

Fast



Memory



Processing

Rich

Spark
SQL

Spark/Stru
ctured
Streaming

Spark
R

Spark
ML

Graph X

Spark Core

Spark Core



Connectors

{ JSON }



JDBC

Parquet



MySQL



amazon web services S3



dBase

APACHE
HBASE

elasticsearch.



amazon web services
Amazon Redshift

and more...



The diagram illustrates the Spark ecosystem architecture. It features a top row of five components: Spark SQL, Spark Streaming, Spark R, Spark ML, and Graph X. Spark SQL is highlighted in blue, while the others are in gray. These components are separated by thin vertical lines. Below this row is a single, wide gray block labeled 'Spark Core', which serves as the foundation for all the components above it.

**Spark
SQL**

**Spark
Streaming**

**Spark
R**

**Spark
ML**

Graph X

Spark Core

Spark SQL

SQL

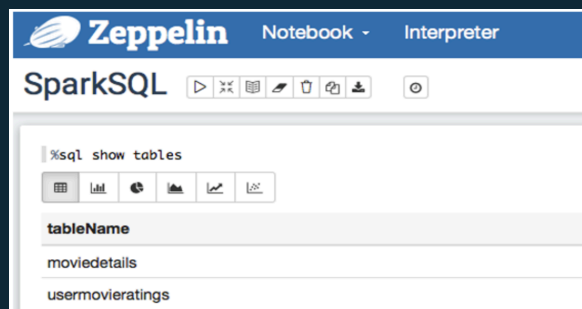


Spark SQL

BI Tool



SQL Editor



CLI

```
select movie-name, avg(rating) r
from moviedetails
group by movie-name
order by r desc
limit 10;
```

Spark
SQL

Spark/Stru
ctured
Streaming


Spark
R


Spark
ML

Graph X

Spark Core

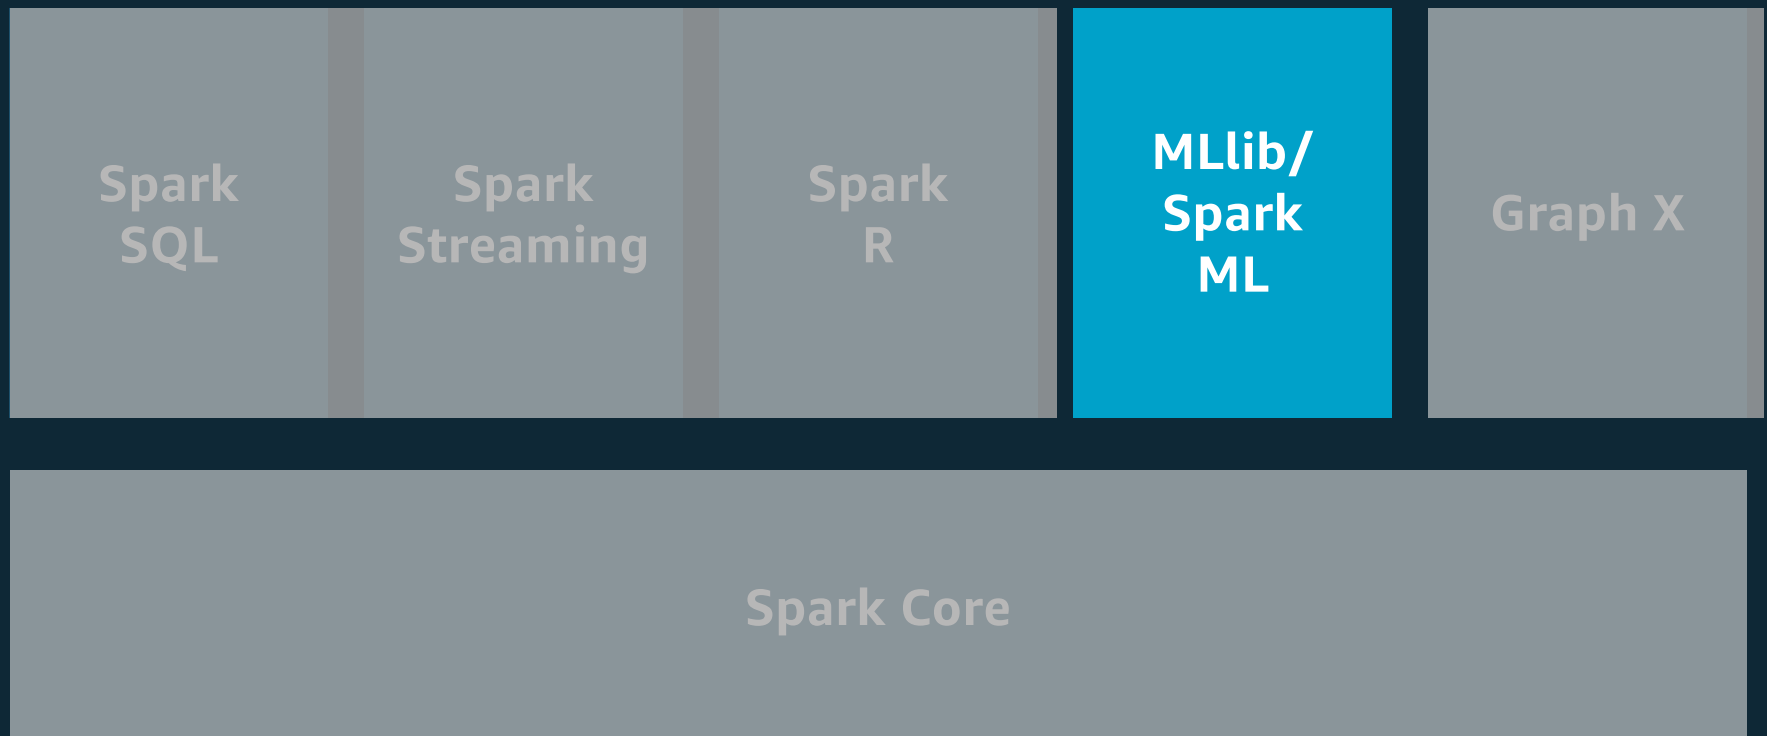
Streaming



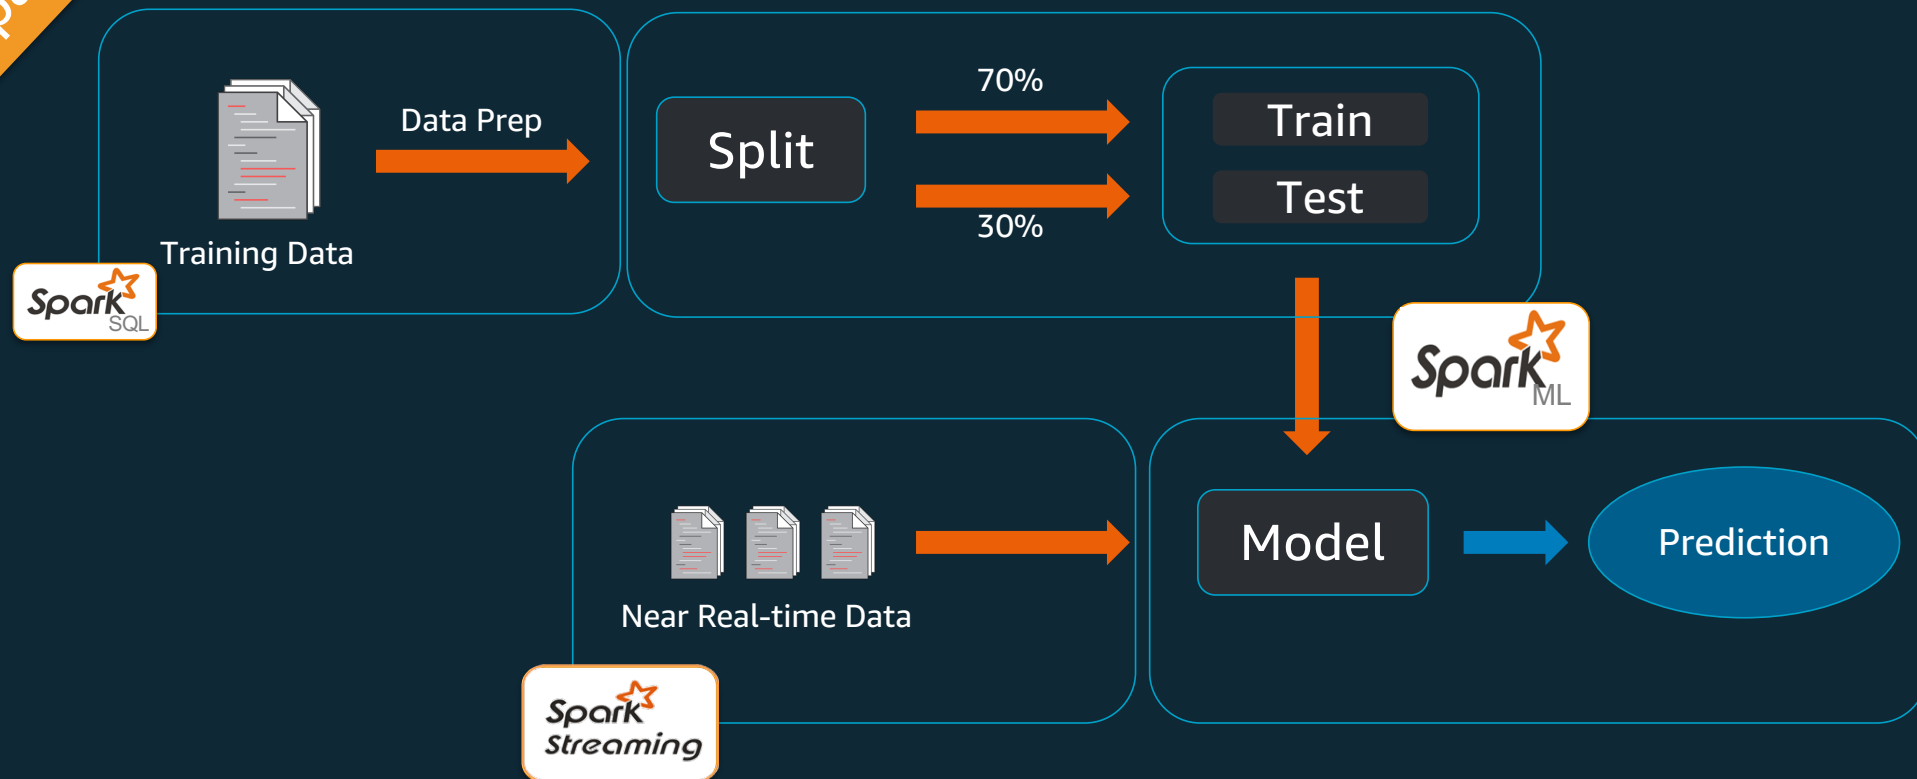



Spark R






Spark ML





Graph X



Apache Spark on Amazon EMR

© 2018, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Launching a cluster with Spark

// CLI

```
aws emr create-cluster --name "Spark cluster" --release-label emr-5.16.0 --applications  
Name=Spark \ --ec2-attributes KeyName=myKey --instance-type m4.large --instance-count 3 --use-  
default-roles
```

// SDK

```
AmazonElasticMapReduceClient emr = new AmazonElasticMapReduceClient(credentials);  
Application sparkApp = new Application().withName("Spark");  
Applications myApps = new Applications();  
myApps.add(sparkApp);  
RunJobFlowRequest request = new RunJobFlowRequest().withName("Spark Cluster")  
.withApplications(myApps).withReleaseLabel("emr-5.16.0").withInstances(new  
JobFlowInstancesConfig().withEc2KeyName("myKeyName").withInstanceCount(1)  
.withKeepJobFlowAliveWhenNoSteps(true).withMasterInstanceType("m4.large")  
.withSlaveInstanceType("m4.large") );  
RunJobFlowResult result = emr.runJobFlow(request);
```

Using Glue Data Catalog as metastore for Spark SQL

- Configure Glue Data Catalog as metastore for Spark SQL (EMR ver > 5.8)
- Persistent metastore – shared by different clusters, services, applications
- Considerations
 - Renaming tables from within Glue – not supported
 - Column statistics, Hive authorization, Hive constraints, Cost based optimization in Hive are not supported
 - Partition values with quotes or apostrophe not supported
 - No UDFs in query predicates
 - Temporary tables not supported

Submitting Spark jobs

- EMR Step API (Console, CLI, SDK)
- `spark-submit` script

Optimizations in Spark

© 2018, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Tricks and tips

- Use newer versions of Spark
- Avoid RDDs
- Cache and reuse results
- Use Parquet format for saving and loading data
- Partition your inputs and outputs
- Avoid shuffling and repartitioning

Tricks and tips

- Pick the right instance types
 - If memory is a constraint, use m4 instances instead of c4 instances
- Configure Spark correctly
 - default EMR configuration for Spark are very simplistic
 - For e.g., enable speculative execution
 - Enable/Disable Dynamic allocation based on workload
 - Minimize number of executor instances per machine
 - Avoid executors with too much memory
 - Number of active tasks = number of cores

Tricks and tips

- Avoid 'LIKE' in SQL queries
- Reduce data shuffles
- Optimize joins
 - Avoid cross joins
 - Use broadcast join for smaller datasets
 - Use flatmap instead of broadcast join
- Avoid spilling to disk
- Detect and handle data skew